

Agent Ops Mission Control

A Compliance Firewall & Governance Platform for AI Agent Fleets

Berlin AI Labs
Gopalkrishna Yamijala · Berlin, Germany
Live Platform · berlinailabs.de

The EU AI Act Takes Effect August 2026

AI systems in healthcare, finance, and insurance are classified high-risk — facing fines up to €35M for non-compliance. Enterprises deploying autonomous AI agents need runtime governance: not monitoring, not logs, but real-time enforcement. No production-ready solution exists today.

€35M

MAX FINE PER VIOLATION

Aug '26

ENFORCEMENT DEADLINE

0

PRODUCTION SOLUTIONS TODAY

WHAT WE BUILT

Mission Control is a **real-time compliance firewall** that enforces regulatory rules on every AI agent conversation — preventing violations before they happen.

- ▶ **GDPR Enforcement** — 8 articles enforced with working code: consent detection, crisis intervention, data residency, transparency
- ▶ **Local AI Firewall** — Checks every conversation in 8ms median latency. No data leaves the device. No cloud dependency
- ▶ **Fast Audit** — Paste any AI transcript, get a compliance report with regulator-ready XML documentation in seconds
- ▶ **Containment** — Circuit breakers halt rogue agents instantly. Kill switches enforce human override
- ▶ **Security Testing** — Open-source CLI on npm (agent-pentest) runs 41 adversarial attack vectors against any AI endpoint
- ▶ **Integration Exports** — REST API, Webhooks, BfArM XML, SIEM (Splunk/Datadog), PDF Audit Reports

⚡ Independent Innovation — Validated by Industry

We independently built the Agent Trust Protocol and PDP Protocol (Proof of Delivery) **before** the Cloud Security Alliance published their Agentic Trust Framework. When their specification was released, it confirmed the architecture we had already deployed:

Agent Identity Behavioral Auditing Data Governance

Access Segmentation Incident Response

BENCHMARKS (PUBLICLY VERIFIABLE)

97%

CRISIS DETECTION (F1:
0.97, N=312)

96%

CONSENT PRECISION
(F1: 0.94, N=847)

8ms

MEDIAN LATENCY
(LOCAL)

12

DEPLOYED
MICROSERVICES

All benchmarks verifiable at </api/gdpr/benchmarks>. Known limitations disclosed on public transparency page.

AGENT GOVERNANCE

- ▶ **Reputation Registry** — Every agent gets a verified compliance score based on audits and behavioral analysis
- ▶ **Maturity Model** — 4-level progression (Intern → Junior → Senior → Principal) with automated promotion gates
- ▶ **Cryptographic Proofs** — Immutable audit trail anchoring every agent action for regulatory non-repudiation

VALIDATION

✓ CIC Berlin Demo

✓ HelloBetter Founder Feedback

✓ Open Source on GitHub

✓ agent-pentest on npm

✓ Live on Railway

WHY RISING INNOVATORS

Every enterprise deploying AI agents in Europe will need runtime governance by August 2026. The market has published specifications — **we shipped production code first**. Zürich recognition connects this infrastructure to the enterprises and partners who need it most.